



The AI grants team for small nonprofits.

Finds the federal grants you can actually win · tells you honestly which to skip
· drafts reviewer-critiqued proposals in minutes.

Shivam Gupta · AI Builders Hackathon 2026 · Live Grants.gov data + a team of Claude agents

THE PROBLEM

The money exists. The capacity doesn't.

Small nonprofits — food pantries, youth programs, shelters — are the delivery layer of American social policy. Federal grants are written for them. Yet the application system assumes staff they don't have.



1.8M

US nonprofits — the overwhelming majority with no development staff or grant writer



\$80-150/hr

market rate for professional grant writers, out of reach for small budgets



40+ hrs

of staff time per federal application — a month of an ED's spare capacity

The result: billions in grants go unclaimed or flow to whoever could afford the writer — while hundreds of opportunities sit open on Grants.gov *right now* (our landing page shows the live count).

Nonprofits don't lose grants because they write badly.

They lose because they apply to the wrong grants — and run out of time.

A wrong application costs a small org a month of capacity. Every competitor builds a writing tool. Granted starts where the money is actually lost:

Eligibility as a hard gate

funder's applicant-type rules,
checked first

Expected-value math, shown

$\text{award} \times \text{win odds} - \text{cost to apply}$

An honest “skip this one”

the AI that says no is the one you
trust

THE PRODUCT

Mission statement in. Reviewed federal draft out.



1 • Profile

Paste anything — mission, website copy. Claude extracts a structured profile of real programs and outcomes. Extraction only; it never invents your numbers.



2 • Triage

Live Grants.gov search + a go/no-go brief per grant: eligibility gate, alignment, capacity vs. award size, effort, and the EV math — including honest skips.




3 • Draft

A streaming agent team writes the proposal, scores it like the funder's panel, and revises what got flagged — before you ever see it.



4 • Track

A deadline-sorted pipeline with the combined expected value of everything you're chasing.



Live product — the counter on the hero is the real number of open federal opportunities.

IT WORKS — LIVE

Watch a real \$1M opportunity go from fit brief to reviewed draft.

The screenshot shows the 'granted.' interface with a navigation bar including 'Discover', 'Pipeline 1', and 'Organization'. The main content area is titled 'Brightpath Youth Alliance' and is in 'Demo mode'. It displays a 'WHY IT FITS' section with three checkmarks: 'Strong keyword overlap between your focus areas and the funder's synopsis.', 'The opportunity sits in a funding category you selected.', and 'Award size is well matched to your organizational budget.'. Below this is an 'ELIGIBILITY' section with the status 'Eligible' and a note: 'Your organization type matches the funder's listed applicant types.'. To the right is a 'RISKS & GAPS' section with a warning: '! Only 9 days to the deadline — very tight for a federal application.'. Further right is an 'EFFORT TO APPLY' section with '~40 Hrs • Medium' and a description: 'Typical federal narrative + budget + registration workload for an opportunity of this size (heuristic estimate)'. To the right of that is an 'EXPECTED VALUE' section with '+\$56,000' and a note: 'Assumes a ~12% win rate for a small applicant and ~\$4,000 of staff/consultant time to apply.'. Below these is an 'IF YOU APPLY — HOW TO WIN' section with three numbered steps: '1. Lead with your most quantified outcome — federal reviewers score evidence, not adjectives.', '2. Mirror the funder's own vocabulary from the synopsis in your need statement.', and '3. Name specific local partners to answer the capacity question before it's asked.'. At the bottom is a 'Drafting complete' section with buttons for 'Redraft from scratch', 'Download .md', and 'Copy'. It shows a progress bar with four agents: 'Strategist' (Planned 7 sections around the funder's priorities.), 'Writer' (First draft complete.), 'Reviewer' (Panel score: 74/100. Strong model and evidence, but the need statement won't survive a BJA panel as written.), and 'Reviser' (Revision complete. Updated panel score: 88/100.). Below the progress bar is a 'STRATEGIST'S READ ON THIS FUNDER' section with the text: 'BJA is buying three things in this notice: credentialed vocational training aligned to the local job market, measurable job placement and retention, and'.

74 → 88

panel score before/after the Reviser addresses the Reviewer's critical notes — with per-section diffs

“Skip.”


a perfect-sounding grant rejected because its eligibility list excludes nonprofits — one honest no saves 40 hours

≈ 3 min

from click to reviewed, exportable draft — streamed live so you watch every agent work

BJA Second Chance Act (real opportunity, up to \$1M) — fit brief with EV, then the four-agent run.

Judgment and generation are separate agents.



- Claude Opus 5 end to end (model-swappable via one env var)
- Structured outputs (Zod schemas) for every judgment — the UI never parses prose
- One streaming call writes all sections; a chunk-safe marker parser routes them live (unit-tested to 1-char chunks)
- Deterministic code owns the arithmetic — EV, capacity ratios, eligibility codes. The model owns judgment.

BUILT FOR TRUST

AI a grant reviewer — and your board — can trust.



It says no.

Wrong applications are the enemy, not the revenue. Ineligible or bad-EV grants get a reasoned “skip” — before you spend a minute.



[ADD: your data]

The Writer never invents statistics, partners, or outcomes. Missing facts become visible placeholders — in federal grants, a fabricated number is disqualifying.



It shows the math.

Every recommendation carries its reasoning and economics: realistic award × conservative win odds — true cost to apply. The decision stays yours.

Every fallback is labeled in the UI: heuristic engine without a key, snapshot data if Grants.gov is down, typed errors from Claude — the product never pretends.

WHY THIS IS A BUSINESS

Software margins on a \$3,000 pain.

What one reviewed federal proposal costs

Professional grant writer



Staff DIY (40+ hrs of an ED's time)



Granted (Claude Opus 5 compute)



Fit brief: ≈ \$0.05-0.08 · a heavy user drafting 10 proposals/mo costs ≈ \$4.50 to serve

TARGET USERS

Small/mid nonprofits (\$100K-\$5M budgets) + grant consultants who serve many orgs at once

PRICING

Starter \$49/mo · Team \$149/mo — vs. \$1.5-3K/yr discovery-only databases and \$3-5K/proposal writers


WEDGE → MOAT

Free eligibility triage on live data → the org's outcome profile + win/loss history compounds

EXPANSION

Foundation & state sources, compliance reporting, consultant multi-org seats

Boring where it should be. Tested where it hurts.



Every dependency has a tested fallback

No API key	transparent heuristic fit engine + replayed pipeline demo — labeled, never a wall
Grants.gov down	bundled snapshot of real opportunities, labeled “offline”
Claude errors / refusal / token cap	typed errors → friendly retry; SSE emits a terminal error event
Malformed stream frames	skipped, never crash the draft; parser tested to 1-char chunks

26

unit tests + a Playwright E2E that drives the whole journey — it caught a real state-loss bug before any human did

WHAT'S NEXT

From federal MVP to the grants OS for small nonprofits.

Now

Federal grants, live · agent pipeline · honest triage (this demo)

Next 3 mo

Foundation & state sources (990-PF) · teams (Postgres + auth) · DOCX/JustGrants export

6-12 mo

Win/loss feedback → calibrated EV priors · funder-specific Reviewer rubrics (BJA, DOL, HHS)

The moat

The compounding dataset of org outcomes × funder decisions — every proposal makes the next one smarter

The grant your community needs is open right now.

github.com/shi1720/ai-b-h · works with zero keys · Shivam Gupta — AI Builders Hackathon 2026